

# ECE465/565 – Operating Systems Design: Project #1

Due date: September 17, 2023

## Objectives

- To become familiar with the general structure of Xinu codebase;
- To become familiar with Xinu's process management;
- To practice writing and understanding test case files;
- To understand the handling of the stack.

## Reminders on Code Reuse, Testing and Use of ChatGPT.

**Code reuse:** This course has a **strict no-code reuse policy**: *consulting or reusing* (even partially) code from other individuals or online resources is not allowed (and is considered an academic integrity violation).

**Testing:** Your code will be graded using the course's VCL image. Please make sure you test your code on the VCL before submitting.

**ChatGPT use:** The use of ChatGPT (<https://chat.openai.com>) is allowed with the restrictions indicated in the syllabus (please consult the syllabus for guidelines). If you use ChatGPT, please list all the questions entered in it in the following form: <https://forms.gle/HgGTSVzF8CynFiCq5>.

## Getting Oriented in Xinu

### Types and Constants

All header files are in the `include` directory. You may want to start from the following header files:

- `include/xinu.h`: unifies the inclusion of all necessary header files.
- `include/prototypes.h`: declares most system-call prototypes.
- `include/kernel.h`: contains definition of some important constants, types, and function prototypes.

### For Debugging

See `kprintf()` system call in `system/kprintf.c`

### Process-related code

- Process definition and PCB table: `include/process.h`
- Process creation/termination/suspension/resumption: `system/create.c`, `system/kill.c`, `system/exit.c`, `system/suspend.c`, `system/resume.c`

- Process scheduling: `include/resched.h`, `system/resched.c`
- Context switch and state change: `system/ctxsw.S`, `system/ready.c`, `system/sleep.c`, `system/unsleep.c`, `system/wait.c`, `system/yield.c`, `system/wakeup.c`
- Process queue management: `include/queue.h`, `system/queue.c`, `system/newqueue.c`, `system/getitem.c`, `system/insert.c`

## Bootstrap procedure

The startup code (`system/start.S`) invokes `nulluser()` in `system/initialize.c` to initialize the system. Analyze the initialization code in `initialize.c`.

## Questions (include the answers to these questions in the report):

- Q1. What is the maximum number of processes accepted by Xinu? Where is it defined?
- Q2. What does Xinu define as an “illegal PID”?
- Q3. What is the default stack size Xinu assigns each process? Where is it defined?
- Q4. Draw Xinu’s process state diagram.
- Q5. When is the shell process created?
- Q6. Draw Xinu’s process tree (including the name and identifier of each process) when the initialization is complete.
- Q7. (Please see Q7 below)

## Coding problems

### P1: Cascading termination

The goal of this problem is to implement cascading termination. For this problem, we will group processes in two categories: *system processes* and *user processes*. We will call “system processes” the processes spawned by default by Xinu (e.g., `startup`, `main`, etc.), and “user processes” the ones spawned by the `main` process (and used to model user code). You can consider the `shell` to be a system process. To distinguish between the two categories of processes, add a `user_process` flag, of type `bool8`, to the PCB, and set the value of this flag to `TRUE` for user processes, and to `FALSE` for system processes. You can modify the `create` function to set this flag to `FALSE` by default, and later modify it as needed.

1. Modify Xinu to support cascading termination *for user processes only*.
2. Write a test case file to validate your implementation. Your test case file should:
  - a. Spawn at least 15 user processes using the `create` system call. The structure of the process tree is up to you, but should be designed to allow for testing representative scenarios (e.g., termination of a leaf process, etc.)
  - b. Kill select processes to test representative scenarios.
  - c. Print the list of active processes before and after process termination. For each active process, print the identifiers of its children.
  - d. Do **not** use the `shell` process in your test cases.

In your submission, the test case file should be named `main.kill`. Include in the report an illustration of the process tree used to test your code. If you don't know how to structure the test case file, have you look at `main.fork` as a reference. You will use `main.fork` to test your code for P2.

## P2: Process creation and stack handling

Implement a `fork` system call similar to Unix's `fork`. The implementation should be in a separate `fork.c` file within the `system` folder. The function declaration must be as follows:

```
pid32 fork()
```

The `fork` primitive creates a new process (the child) by *almost* duplicating the parent process. Note that Xinu's processes are essentially threads: they share the memory address space but have *private stacks*.

**Initialization:** The child process should be initialized to have the same name, priority and stack length as the parent process. The `prsem`, `prhasmg` and `prdesc` fields of its process control block (PCB) should be initialized as during standard process creation (to `-1`, `FALSE` and `CONSOLE`, respectively). The child process should be set in READY state (and inserted in the ready list) upon creation. The ready list contains the processes that are eligible for execution. Have a look at `ready.c` (line 25) to see how to add a process to the ready list.

**Return value:** On success, `fork` must return the PID of the child process to the parent, and value `NPROC` to the child (note: differently from Unix's `fork`, the function you are coding does not return value 0 to the child because that's a PID in use). On failure, `fork` returns `SYSERR`.

**Execution:** After creation, *the child process should resume execution starting from the first instruction following the `fork` system call.*

**Testing:** Test cases for the `fork` system call are provided in `main.fork`. To run these test cases, rename `main.fork` into `main.c`, and compile the code. Note that you can selectively disable the test cases by commenting out the respective `"#define TESTCASEx"` macro definition at the beginning of the file. The expected output is contained in the provided `fork.output` file.

**Hints:** This problem requires a good understanding of the stack and its handling by Xinu. To this end:

- You can find information on the x86 architecture and assembly in the Intel Architecture Software Developer's Manual (available on the course website). Refer to Volume 1, Chapters 4.1-4.3 for information on the handling of the stack. You can also find a brief introduction on the runtime stack in the Xinu's book, Chapter 3.9.1.
- Have a look at `create.c` and `ctxsw.S` to see how the stack is initialized and how it is handled upon context switch.
- The `stacktrace` system call (in `stacktrace.c`) prints the stack backtrace for a given process, and can help you understand how to traverse the stack.
- Have a look at the three test cases in `main.fork` and make sure that you fully understand their expected output (in `fork.output`).
- If you want to see assembly file corresponding to any Xinu file, you can proceed as follow. First, modify the `Makefile` (in the `compile` folder) and add `"-S"` to the `CFLAGS` compiler flag. Second, compile using `"make clean; make;"`. After compilation, the `.o` files in the `binaries` subfolder will contain assembly (rather than object) code.

**Question Q7:** what is the effect of the “`receive()`” call in the test cases provided? What would happen if that function call was not present? *Hint:* see `receive.c` and `send.c` files. Include the answer to this question in the report.

### Submission instructions

1. Create a `testcases` folder in the `system` folder and include the test case related files (`main.fork`, `fork.output` and `main.kill`) in it.
2. **Important:** You can write code in `main.c` to test your functions, but please note that, when we test your code, we may replace the `main.c` file. Therefore, do not implement any essential functionality in the `main.c` file. Also, turn off debugging output before submitting your code.
3. Go to the `xinu/compile` directory and invoke `make clean`.
4. Create a `xinu/tmp` directory and copy all the files you have modified/created (both `.h` files and `.c` files) into it (the `tmp` folder should have the same directory structure as the `xinu` folder). For example, if you have modified `Makefile`, `system/create.c` and `open.c`, your `xinu` directory will look like:

```
-xinu
  -[existing folders]
  -tmp [this is the folder you created]
    -compile
      Makefile
    -system
      create.c open.c
    - testcases
      main.kill
```

Note that the use of the `tmp` folder aims to help the TA to quickly identify what files have been modified. Please **do not delete** the files that you have created/modified from the original folders (i.e., the `xinu/include`, `xinu/system`, etc., you have been working on).

5. The project report should contain:
  - The answer to all questions posed in this assignment (Q1-Q7);
  - For each coding problem, a **brief** description of your implementation (which files have you added/modified? What are the main data structures used by your implementation?);
  - For P1, the illustration of the process tree used to test your code.
6. Go to the parent folder of the `xinu` folder. Compress the whole `xinu` directory into a `tgz` file.  
`tar czf xinu_project1.tgz xinu`
7. Submit your code and report (in pdf format) through Moodle. *Please do not include the report in the `xinu_project1.tgz` archive; submit code and report as two separate files.*